

MultiLayered Neural Network

Objective : Wokring with Keras and Experiment with Different Nueral Network Architectures.

Dataset : Mnist DataSet which is avialalbe in Keras Dataset by default.

Developer Details : PraveenAI

Source Details : Most of the code is extracted from <https://keras.io/losses/> (<https://keras.io/losses/>) and Base Archictural Refference is from <https://github.com/wagonhelm/NaNmnist/blob/master/NaNmnist.ipynb> (<https://github.com/wagonhelm/NaNmnist/blob/master/NaNmnist.ipynb>)
https://github.com/aymericdamien/TensorFlow-Examples/blob/master/notebooks/3_NeuralNetworks/neural_network.ipynb (https://github.com/aymericdamien/TensorFlow-Examples/blob/master/notebooks/3_NeuralNetworks/neural_network.ipynb).

Keras

`datasets.mnist.load_data()` :

Can load the Mnist datasets which is available by default provided by Keras.

`np_utils.to_categorical` :

This will help to Convert a class vector (integers) to binary class matrix.(like one-hot coding)

`Sequential()`

To start initializing the linear stack of layers. we can create a Sequential model by passing a list of layer instances to the constructor:

`add(Dense())`

This means we are Adding a Layer to the Sequence.

`input_dim` = this is only for the first layer, so that NN will understand the Input dimensions

`activation`= We provide the Activation function inside the Neurons

`input_shape` = No of neurons inside is depends on the input parameters to the Dense

`kernel_initializer` = helps to Initialize the weights for each connection from one layer to another layer.

Dropout

Fraction of the input units to drop

`BatchNormalization()`

If we have too many layers, there might be a vanishing gradients problem.

To avoid to some extent, Batch Normalization will normalize the activations of the previous layer at each batch.

It will make the mean to 0 and Std Deviation to 1.

`compile()`

This compiles the Sequence that we defined with accepting some input params

`optimizer` = This helps in controlling the Learning Rate of the Weights.

`loss`= we have to input loss function based on the type of that problemset we are working on.

`metrics`= we have to input the Performance metric that we are looking for

In [0]:

```
from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal
from warnings import catch_warnings
```

In [0]:

```
(X_train, y_train), (X_test, y_test) = mnist.load_data()
print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d, %d)"%(X_train.shape[1], X_train.shape[2]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d, %d)"%(X_test.shape[1], X_test.shape[2]))
```

Number of training examples : 60000 and each image is of shape (28, 28)

Number of training examples : 10000 and each image is of shape (28, 28)

Reshape the DataSet Dimensinality

In [0]:

```
# if you observe the input shape its 2 dimensional vector
# for each image we have a (28*28) vector
# we will convert the (28*28) vector into single dimensional vector of 1 * 784
```

```
X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

```
# after converting the input images from 3d to 2d vectors
print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d)"%(X_train.shape[1]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d)"%(X_test.shape[1]))
```

Number of training examples : 60000 and each image is of shape (784)

Number of training examples : 10000 and each image is of shape (784)

Data Normlization

In [0]:

```
# if we observe the above matrix each cell is having a value between 0-255
# before we move to apply machine learning algorithms lets try to normalize the data
# X => (X - Xmin)/(Xmax-Xmin) = X/255
```

```
X_train = X_train/255
X_test = X_test/255
```

Convert Labels to Categories for Softmax Activation(One-Hot coding)

In [0]:

```
# here we are having a class number for each image
print("Class label of first image :", y_train[0])

# Lets convert this into a 10 dimensional vector
# ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
# this conversion needed for MLPs

Y_train = np_utils.to_categorical(y_train, 10)
Y_test = np_utils.to_categorical(y_test, 10)

print("After converting the output into a vector : ",Y_train[0])
```

Class label of first image : 5

After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]

Plot Train Los and Test Loss Vs Epochs

In [0]:

```
%matplotlib inline
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
import matplotlib.pyplot as plt
import numpy as np
import time

def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()
```

In [0]:

```
from keras.models import Sequential
from keras.layers import Dense, Activation
Total_SUMRY = {}
```

Below we are gonna Experiment with three Different Architecture on Mnist Dataset.

We are consdiering Accuracy as performace metric.

Architecture 1 : 784-784-424-10 (2 Hidden Layer with 784,424 Relu)

MLP + Relu activation + AdamOtpimizer

In [0]:

```
#model parameters

batch_size = 128
nb_epoch = 20

input_dim = X_train.shape[1] #784
H_Layer_1 = 784
H_Layer_2 = 424
output_dim = 10
```

In [0]:

```
model_1 = Sequential()
model_1.add(Dense(input_dim ,activation='relu',input_shape=(input_dim,)))
model_1.add(Dropout(0.1))
model_1.add(Dense(H_Layer_1, activation='relu'))
model_1.add(Dropout(0.3))
model_1.add(BatchNormalization())
model_1.add(Dense(H_Layer_2, activation='relu'))
model_1.add(Dropout(0.5))
model_1.add(BatchNormalization())
model_1.add(Dense(output_dim, activation='softmax'))
model_1.summary()
```

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_47 (Dense)	(None, 784)	615440
dropout_28 (Dropout)	(None, 784)	0
dense_48 (Dense)	(None, 784)	615440
dropout_29 (Dropout)	(None, 784)	0
batch_normalization_24 (Batc	(None, 784)	3136
dense_49 (Dense)	(None, 424)	332840
dropout_30 (Dropout)	(None, 424)	0
batch_normalization_25 (Batc	(None, 424)	1696
dense_50 (Dense)	(None, 10)	4250
=====	=====	=====
Total params: 1,572,802		
Trainable params: 1,570,386		
Non-trainable params: 2,416		

We have initialized our model with batch_size is 128 and 20 no of epoch

1st Layer is with 784 Relu activation Function and the 2nd Layer with 424

In [0]:

```
model_1.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy']  
)  
history = model_1.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose  
=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 26s 438us/step - loss: 0.27

31 - acc: 0.9158 - val_loss: 0.1306 - val_acc: 0.9600

Epoch 2/20

60000/60000 [=====] - 23s 390us/step - loss: 0.11

86 - acc: 0.9644 - val_loss: 0.0902 - val_acc: 0.9719

Epoch 3/20

60000/60000 [=====] - 23s 389us/step - loss: 0.08

56 - acc: 0.9735 - val_loss: 0.0855 - val_acc: 0.9734

Epoch 4/20

60000/60000 [=====] - 23s 390us/step - loss: 0.06

65 - acc: 0.9792 - val_loss: 0.0725 - val_acc: 0.9769

Epoch 5/20

60000/60000 [=====] - 23s 390us/step - loss: 0.05

52 - acc: 0.9824 - val_loss: 0.0635 - val_acc: 0.9799

Epoch 6/20

60000/60000 [=====] - 24s 392us/step - loss: 0.04

73 - acc: 0.9853 - val_loss: 0.0730 - val_acc: 0.9776

Epoch 7/20

60000/60000 [=====] - 23s 388us/step - loss: 0.04

22 - acc: 0.9864 - val_loss: 0.0709 - val_acc: 0.9802

Epoch 8/20

60000/60000 [=====] - 23s 389us/step - loss: 0.03

65 - acc: 0.9884 - val_loss: 0.0653 - val_acc: 0.9813

Epoch 9/20

60000/60000 [=====] - 23s 383us/step - loss: 0.03

43 - acc: 0.9883 - val_loss: 0.0732 - val_acc: 0.9789

Epoch 10/20

60000/60000 [=====] - 23s 382us/step - loss: 0.02

95 - acc: 0.9903 - val_loss: 0.0623 - val_acc: 0.9832

Epoch 11/20

60000/60000 [=====] - 23s 388us/step - loss: 0.02

87 - acc: 0.9910 - val_loss: 0.0691 - val_acc: 0.9806

Epoch 12/20

60000/60000 [=====] - 23s 387us/step - loss: 0.02

45 - acc: 0.9913 - val_loss: 0.0697 - val_acc: 0.9820

Epoch 13/20

60000/60000 [=====] - 23s 386us/step - loss: 0.02

31 - acc: 0.9927 - val_loss: 0.0760 - val_acc: 0.9807

Epoch 14/20

60000/60000 [=====] - 23s 384us/step - loss: 0.02

31 - acc: 0.9926 - val_loss: 0.0710 - val_acc: 0.9807

Epoch 15/20

60000/60000 [=====] - 24s 397us/step - loss: 0.02

04 - acc: 0.9936 - val_loss: 0.0610 - val_acc: 0.9835

Epoch 16/20

60000/60000 [=====] - 24s 396us/step - loss: 0.01

75 - acc: 0.9947 - val_loss: 0.0776 - val_acc: 0.9814

Epoch 17/20

60000/60000 [=====] - 24s 393us/step - loss: 0.01

95 - acc: 0.9938 - val_loss: 0.0616 - val_acc: 0.9842

Epoch 18/20

60000/60000 [=====] - 24s 395us/step - loss: 0.01

65 - acc: 0.9948 - val_loss: 0.0736 - val_acc: 0.9816

Epoch 19/20

60000/60000 [=====] - 23s 386us/step - loss: 0.01

75 - acc: 0.9941 - val_loss: 0.0636 - val_acc: 0.9847

Epoch 20/20

60000/60000 [=====] - 23s 382us/step - loss: 0.01

51 - acc: 0.9954 - val_loss: 0.0646 - val_acc: 0.9836

In [0]:

```
score1_tst = model_1.evaluate(X_test, Y_test, verbose=0)
print('Test score    :', score1_tst[0])
print('Test accuracy:', score1_tst[1])

score1_trn = model_1.evaluate(X_train, Y_train, verbose=0)
print('Train accuracy:', score1_trn[1])


fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
x = list(range(1,nb_epoch+1))
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

Total_SUMRY[1]=["Model1", "2", "784-424" ,"Relu","Adam","No","No-Initialization",score1_trn[1],score1_tst[1]]
```

Test score : 0.06464198227513315

Test accuracy: 0.9836

Train accuracy: 0.9982666666666666



As we notice, Loos is so high for Train Data at intial Epochs.

Test Loss for lesser than the Validation lost till the first five epochs.

For every Gradient the loss is converging, but form 15th epoch it is stable for both Test and Validation.

At the 4th epoc bot values are same, 0.18.

Model Performance

Test accuracy: 0.9836 and Train accuracy: 0.999

Seems it is Overfitting, since both accuracies are converging, we can not say it.

it is Overfitting, if the Test accuracy is very low than Train accuracy is so high.

Hence, the Above Architecture is Performing very Well.

Architecture 2:

784-524-462-128-10 (3 Hidden Layer with 524,462,128 Relu)

MLP + Relu Activation + adam Optimizer + RandomNormal Initialization

In [0]:

```
#model parameters

batch_size = 200
nb_epoch = 20

input_dim = X_train.shape[1] #784
H_Layer_1 = 524
H_Layer_2 = 462
H_Layer_3 = 128
output_dim = 10
```

In [0]:

```
# https://arxiv.org/pdf/1707.09725.pdf#page=95
# for relu layers
# If we sample weights from a normal distribution  $N(\theta, \sigma)$  we satisfy this condition with  $\sigma = \sqrt{2/(n_i)}$ .
# hi =>  $\sigma = \sqrt{2/(fan\_in)}$ 
# out =>  $\sigma = \sqrt{2/(fan\_in+1)}$ 
model_2 = Sequential()
model_2.add(Dense(input_dim, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.0618, seed=None)))
model_2.add(Dropout(0.1))
model_2.add(Dense(H_Layer_1, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.0633, seed=None)))
model_2.add(Dropout(0.2))
model_2.add(BatchNormalization())
model_2.add(Dense(H_Layer_2, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.0658, seed=None)))
model_2.add(Dropout(0.45))
model_2.add(BatchNormalization())
model_2.add(Dense(H_Layer_2, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)))
model_2.add(Dropout(0.5))
model_2.add(BatchNormalization())
model_2.add(Dense(output_dim, activation='softmax'))
model_2.summary()
```

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_51 (Dense)	(None, 784)	615440
dropout_31 (Dropout)	(None, 784)	0
dense_52 (Dense)	(None, 524)	411340
dropout_32 (Dropout)	(None, 524)	0
batch_normalization_26 (Batch Normalization)	(None, 524)	2096
dense_53 (Dense)	(None, 462)	242550
dropout_33 (Dropout)	(None, 462)	0
batch_normalization_27 (Batch Normalization)	(None, 462)	1848
dense_54 (Dense)	(None, 462)	213906
dropout_34 (Dropout)	(None, 462)	0
batch_normalization_28 (Batch Normalization)	(None, 462)	1848
dense_55 (Dense)	(None, 10)	4630
=====	=====	=====
Total params:	1,493,658	
Trainable params:	1,490,762	
Non-trainable params:	2,896	

In [0]:

```
model_2.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy']  
)  
history = model_2.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose  
=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 23s 385us/step - loss: 0.42
41 - acc: 0.8682 - val_loss: 0.1256 - val_acc: 0.9605

Epoch 2/20

60000/60000 [=====] - 20s 336us/step - loss: 0.15
19 - acc: 0.9547 - val_loss: 0.1005 - val_acc: 0.9679

Epoch 3/20

60000/60000 [=====] - 20s 338us/step - loss: 0.10
45 - acc: 0.9681 - val_loss: 0.0835 - val_acc: 0.9744

Epoch 4/20

60000/60000 [=====] - 20s 338us/step - loss: 0.07
85 - acc: 0.9768 - val_loss: 0.0724 - val_acc: 0.9787

Epoch 5/20

60000/60000 [=====] - 20s 338us/step - loss: 0.06
12 - acc: 0.9811 - val_loss: 0.0766 - val_acc: 0.9786

Epoch 6/20

60000/60000 [=====] - 20s 338us/step - loss: 0.05
04 - acc: 0.9845 - val_loss: 0.0740 - val_acc: 0.9792

Epoch 7/20

60000/60000 [=====] - 20s 338us/step - loss: 0.04
48 - acc: 0.9859 - val_loss: 0.0684 - val_acc: 0.9815

Epoch 8/20

60000/60000 [=====] - 21s 343us/step - loss: 0.04
24 - acc: 0.9870 - val_loss: 0.0694 - val_acc: 0.9811

Epoch 9/20

60000/60000 [=====] - 21s 343us/step - loss: 0.03
43 - acc: 0.9892 - val_loss: 0.0744 - val_acc: 0.9797

Epoch 10/20

60000/60000 [=====] - 21s 343us/step - loss: 0.03
18 - acc: 0.9903 - val_loss: 0.0791 - val_acc: 0.9787

Epoch 11/20

60000/60000 [=====] - 21s 345us/step - loss: 0.02
98 - acc: 0.9903 - val_loss: 0.0632 - val_acc: 0.9820

Epoch 12/20

60000/60000 [=====] - 21s 343us/step - loss: 0.02
52 - acc: 0.9918 - val_loss: 0.0693 - val_acc: 0.9819

Epoch 13/20

60000/60000 [=====] - 20s 340us/step - loss: 0.02
39 - acc: 0.9927 - val_loss: 0.0722 - val_acc: 0.9825

Epoch 14/20

60000/60000 [=====] - 20s 340us/step - loss: 0.02
60 - acc: 0.9917 - val_loss: 0.0682 - val_acc: 0.9822

Epoch 15/20

60000/60000 [=====] - 20s 342us/step - loss: 0.02
17 - acc: 0.9929 - val_loss: 0.0681 - val_acc: 0.9820

Epoch 16/20

60000/60000 [=====] - 20s 340us/step - loss: 0.02
08 - acc: 0.9934 - val_loss: 0.0776 - val_acc: 0.9807

Epoch 17/20

60000/60000 [=====] - 21s 342us/step - loss: 0.01
99 - acc: 0.9940 - val_loss: 0.0722 - val_acc: 0.9812

Epoch 18/20

60000/60000 [=====] - 21s 342us/step - loss: 0.01
71 - acc: 0.9945 - val_loss: 0.0735 - val_acc: 0.9823

Epoch 19/20

60000/60000 [=====] - 20s 341us/step - loss: 0.01
65 - acc: 0.9948 - val_loss: 0.0752 - val_acc: 0.9814

Epoch 20/20

60000/60000 [=====] - 21s 343us/step - loss: 0.01
61 - acc: 0.9948 - val_loss: 0.0781 - val_acc: 0.9834

In [0]:

```
score2_tst = model_2.evaluate(X_test, Y_test, verbose=0)
print('Test score   :', score2_tst[0])
print('Test accuracy:', score2_tst[1])


score2_trn = model_2.evaluate(X_train, Y_train, verbose=0)
print('Train score   :', score2_trn[0])
print('Train accuracy:', score2_trn[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

x = list(range(1,nb_epoch+1))
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

Total_SUMRY[2]=["Model12", "3", "524-462-128" ,"Relu","Adam","Yes","Random Normal",score
2_trn[1],score2_tst[1]]
```

Test score : 0.07809708171830716
 Test accuracy: 0.9834
 Train score : 0.005488861975786434
 Train accuracy: 0.9982



As we notice, Los is so high for Train Data at intial Epochs.

Test Loss for lesser and constan for most of the epoch with very litile changes.

For every Gradient the validation loss is converging. Even form 15th epoch also, it is stable for both Test and Validation.

At the 4th epoc bot values are same, 0.08.

Model Performance

Test accuracy: 0.9834 and Train accuracy: 0.992

Seems it is Overfitting, since both accuracies are with small difference., we can not say it.

It is Overfitting, if the Test accuracy is very low and the Train accuracy is so high.

Hence, the Above Architecture is Performing very Well.

Architecture 3 : 784-500-425-350-124-64-10

MLP + Relu Activation + Adam Optimizer + BatchNormalization + Dropout

In [0]:

```
from keras.layers import Dropout
from keras.layers.normalization import BatchNormalization
```

In [0]:

```
#model parameters

batch_size = 200
nb_epoch = 20

input_dim = X_train.shape[1] #784
H_Layer_1 = 500
H_Layer_2 = 425
H_Layer_3 = 350
H_Layer_4 = 124
H_Layer_5 = 64
output_dim = 10
```

In [0]:

```
# https://arxiv.org/pdf/1707.09725.pdf#page=95
# for relu layers
# If we sample weights from a normal distribution  $N(\theta, \sigma)$  we satisfy this condition with  $\sigma = \sqrt{2/(n_i)}$ .
# hi =>  $\sigma = \sqrt{2/(fan\_in)}$ 
# out =>  $\sigma = \sqrt{2/(fan\_in+1)}$ 

model_3 = Sequential()
model_3.add(Dense(input_dim, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.0633, seed=None)))
model_3.add(Dense(H_Layer_1, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.0632, seed=None)))
model_3.add(Dropout(0.2))
model_3.add(Dense(H_Layer_2, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.0685, seed=None)))
model_3.add(Dropout(0.3))
model_3.add(BatchNormalization())
model_3.add(Dense(H_Layer_3, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.0756, seed=None)))
model_3.add(BatchNormalization())
model_3.add(Dropout(0.4))
model_3.add(BatchNormalization())
model_3.add(Dense(H_Layer_4, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.127, seed=None)))
model_3.add(BatchNormalization())
model_3.add(Dropout(0.45))
model_3.add(Dense(H_Layer_5, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.1768, seed=None)))
model_3.add(BatchNormalization())
model_3.add(Dropout(0.5))
model_3.add(Dense(output_dim, activation='softmax'))
model_3.summary()
```

Layer (type)	Output Shape	Param #
dense_56 (Dense)	(None, 784)	615440
dense_57 (Dense)	(None, 500)	392500
dropout_35 (Dropout)	(None, 500)	0
dense_58 (Dense)	(None, 425)	212925
dropout_36 (Dropout)	(None, 425)	0
batch_normalization_29 (Batch Normalization)	(None, 425)	1700
dense_59 (Dense)	(None, 350)	149100
batch_normalization_30 (Batch Normalization)	(None, 350)	1400
dropout_37 (Dropout)	(None, 350)	0
batch_normalization_31 (Batch Normalization)	(None, 350)	1400
dense_60 (Dense)	(None, 124)	43524
batch_normalization_32 (Batch Normalization)	(None, 124)	496
dropout_38 (Dropout)	(None, 124)	0
dense_61 (Dense)	(None, 64)	8000
batch_normalization_33 (Batch Normalization)	(None, 64)	256
dropout_39 (Dropout)	(None, 64)	0
dense_62 (Dense)	(None, 10)	650
Total params: 1,427,391		
Trainable params: 1,424,765		
Non-trainable params: 2,626		

In [0]:

```
model_3.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy']  
)  
history = model_3.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose  
=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 24s 396us/step - loss: 0.95

11 - acc: 0.7104 - val_loss: 0.2158 - val_acc: 0.9377

Epoch 2/20

60000/60000 [=====] - 19s 322us/step - loss: 0.26

19 - acc: 0.9315 - val_loss: 0.1222 - val_acc: 0.9639

Epoch 3/20

60000/60000 [=====] - 20s 327us/step - loss: 0.16

88 - acc: 0.9578 - val_loss: 0.1105 - val_acc: 0.9697

Epoch 4/20

60000/60000 [=====] - 19s 322us/step - loss: 0.12

61 - acc: 0.9682 - val_loss: 0.1128 - val_acc: 0.9696

Epoch 5/20

60000/60000 [=====] - 20s 326us/step - loss: 0.09

96 - acc: 0.9748 - val_loss: 0.0972 - val_acc: 0.9773

Epoch 6/20

60000/60000 [=====] - 19s 324us/step - loss: 0.08

18 - acc: 0.9797 - val_loss: 0.0980 - val_acc: 0.9766

Epoch 7/20

60000/60000 [=====] - 19s 320us/step - loss: 0.07

35 - acc: 0.9816 - val_loss: 0.0935 - val_acc: 0.9770

Epoch 8/20

60000/60000 [=====] - 19s 319us/step - loss: 0.06

02 - acc: 0.9848 - val_loss: 0.1078 - val_acc: 0.9755

Epoch 9/20

60000/60000 [=====] - 19s 322us/step - loss: 0.05

40 - acc: 0.9859 - val_loss: 0.0978 - val_acc: 0.9777

Epoch 10/20

60000/60000 [=====] - 19s 322us/step - loss: 0.04

84 - acc: 0.9884 - val_loss: 0.0941 - val_acc: 0.9800

Epoch 11/20

60000/60000 [=====] - 19s 321us/step - loss: 0.04

61 - acc: 0.9886 - val_loss: 0.0970 - val_acc: 0.9792

Epoch 12/20

60000/60000 [=====] - 19s 321us/step - loss: 0.03

64 - acc: 0.9903 - val_loss: 0.0936 - val_acc: 0.9804

Epoch 13/20

60000/60000 [=====] - 19s 324us/step - loss: 0.03

12 - acc: 0.9919 - val_loss: 0.0931 - val_acc: 0.9812

Epoch 14/20

60000/60000 [=====] - 19s 323us/step - loss: 0.03

17 - acc: 0.9920 - val_loss: 0.1015 - val_acc: 0.9809

Epoch 15/20

60000/60000 [=====] - 19s 321us/step - loss: 0.03

29 - acc: 0.9918 - val_loss: 0.1329 - val_acc: 0.9752

Epoch 16/20

60000/60000 [=====] - 19s 322us/step - loss: 0.02

95 - acc: 0.9929 - val_loss: 0.1078 - val_acc: 0.9775

Epoch 17/20

60000/60000 [=====] - 19s 320us/step - loss: 0.02

59 - acc: 0.9938 - val_loss: 0.1073 - val_acc: 0.9795

Epoch 18/20

60000/60000 [=====] - 19s 320us/step - loss: 0.02

43 - acc: 0.9940 - val_loss: 0.0983 - val_acc: 0.9823

Epoch 19/20

60000/60000 [=====] - 20s 325us/step - loss: 0.02

55 - acc: 0.9936 - val_loss: 0.0820 - val_acc: 0.9832

Epoch 20/20

60000/60000 [=====] - 19s 323us/step - loss: 0.01

93 - acc: 0.9950 - val_loss: 0.0890 - val_acc: 0.9836

In [0]:


```
score3_tst = model_3.evaluate(X_test, Y_test, verbose=0)
print('Test score   :', score3_tst[0])
print('Test accuracy:', score3_tst[1])

score3_trn = model_3.evaluate(X_train, Y_train, verbose=0)
print('Train score   :', score3_trn[0])
print('Train accuracy:', score3_trn[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

x = list(range(1, nb_epoch+1))
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
Total_SUMRY[3]=["Model3", "5", "500-425-350-124-64" , "Relu", "Adam", "Yes", "Random Norma
l", score3_trn[1], score3_tst[1]]
```

Test score : 0.08903393339179748
 Test accuracy: 0.9836
 Train score : 0.010572065364455194
 Train accuracy: 0.9974



This is almost similar to the first Architecture

At the 4th epoc bot values are same, 0.17.

As we notice, Los is so high for Train Data at intial Epochs.

Test Loss for lesser and constan for most of the epoch with very litile changes.

At the 4th epoc bot values are same, 0.17.

Model Performance

Test accuracy: 0.9836 and Train accuracy: 0.9974

Seems it is Overfitting, since both accuracies are with small difference., we can not say it.

It is Overfitting, if the Test accuracy is very low and the Train accuracy is so high.

Hence, the Above Architecture is Performing very Well.

In [0]:

```
#http://zetcode.com/python/prettytable/
from prettytable import PrettyTable
x = PrettyTable()
x.field_names = ["Model", "No of Layers", "No of Nuerons", "Activation", "Optimizer", "Dropout", "Initialization", "Test ACC", "Train ACC"]
for i,j in enumerate(Total_SUMRY):
    x.add_row([Total_SUMRY[j][0], Total_SUMRY[j][1], Total_SUMRY[j][2], Total_SUMRY[j][3], Total_SUMRY[j][4], Total_SUMRY[j][5], Total_SUMRY[j][6], Total_SUMRY[j][7], Total_SUMRY[j][8]])
    #print(Total_SUMRY[j][0], " = ", Total_SUMRY[j][1], " = ", Total_SUMRY[j][2], " = ", Total_SUMRY[j][3], " = ", Total_SUMRY[j][4], " = ", Total_SUMRY[j][5], " = ", Total_SUMRY[j][6], " = ", Total_SUMRY[j][7])
print(x)
```

Model	No of Layers	No of Nuerons	Activation	Optimizer	Dropout	Initialization	Test ACC	Train ACC
Model1	2	784-424	Relu	Adam		No-Initialization	0.9982666666666666	0.9836
Model2	3	524-462-128	Relu	Adam		Random Normal	0.9982	0.9834
Model3	5	500-425-350-124-64	Relu	Adam		Random Normal	0.9974	0.9836

Conclusion

We are getting almost all nearest accuracy values. But still out of three Model1 is performing little bit higher.